



## **Lab 3: Linear Regression with a Single Variable**

### **Objective:**

To understand and implement Linear Regression with a single independent variable. The lab focuses on the core concepts, assumptions, and practical implementation using Python. Students will learn to visualize the relationship, interpret results, and evaluate model performance.

### **Prerequisites:**

- Knowledge of basic Python programming.
- Familiarity with libraries: Pandas, NumPy, Matplotlib, and Scikit-learn.
- Understanding of basic statistical concepts like correlation and variance.

### **1. Introduction to Regression Analysis**

#### **What is Regression?**

Regression is a statistical method for understanding the relationship between a dependent variable (target) and one or more independent variables (predictors).

In **Simple Linear Regression**, the relationship is modeled as a straight line:

$$y = \beta_0 + \beta_1 x + \epsilon$$

- $y$ : Dependent variable (e.g., Price).
- $x$ : Independent variable (e.g., Avg. Area House Age).
- $\beta_0$ : Intercept (value of  $y$  when  $x$  is 0).
- $\beta_1$ : Slope (rate of change in  $y$  with a unit increase in  $x$ ).
- $\epsilon$ : Error term (difference between predicted and actual values).



## 2. Assumptions of Linear Regression

1. **Linearity:** The relationship between xxx and yyy is linear.
2. **Independence:** Observations are independent of each other.
3. **Homoscedasticity:** Residuals (errors) have constant variance.
4. **Normality of Residuals:** Residuals follow a normal distribution.

## 3. Theory: Evaluating Linear Regression Models

### Key Metrics:

1. **Mean Squared Error (MSE):** Measures average squared error between predicted and actual values.

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

2. **R-squared ( $R^2$ ):** Explains the proportion of variance in y explained by x

$$R^2 = 1 - \frac{\text{Sum of Squared Errors}}{\text{Total Sum of Squares}}$$

### 3. Visualization:

- Regression line: Shows how well the model fits the data.
- Residual plot: Helps assess assumptions of linearity and homoscedasticity.

## 4. Implementation of Simple Linear Regression

### Step 1: Import Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
```



- **Pandas:** For loading and manipulating tabular data.
- **NumPy:** For mathematical operations (used for arrays and calculations).
- **Matplotlib:** For creating visualizations like scatter plots and regression lines.
- **Scikit-learn:** For building and evaluating the regression model.

### Step 2: Load Dataset

```
data = pd.read_csv("housing.csv")  
print("First 5 rows of the dataset:")  
print(data.head())
```

- `pd.read_csv`: Reads a CSV file into a DataFrame.
- `data.head()`: Displays the first five rows to check data loading and structure.

### Step 3: Data Preprocessing

1. Select the independent (x) and dependent (y) variables.
2. Handle missing values (if any).

```
# Selecting independent (X) and dependent (y) variables  
X = data[['Avg. Area House Age']] # Independent variable  
y = data['Price'] # Dependent variable
```

```
# Check for missing values  
print("\nMissing values in the dataset:")  
print(data.isnull().sum())
```

- `X = data[['SquareFeet']]`: Extracts the independent variable column (predictor).
- `y = data['Price']`: Extracts the dependent variable column (target).
- `isnull().sum()`: Identifies missing values in each column.



#### Step 4: Train-Test Split

Divide the data into training and testing sets.

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

**train\_test\_split:** Divides the data into training and testing sets.

- **test\_size=0.2:** Reserves 20% of the data for testing.
- **random\_state=42:** Ensures reproducibility of the split.

#### Step 5: Train the Model

```
# Initialize and train the linear regression model
model = LinearRegression()
model.fit(X_train, y_train)
```

```
# Print model parameters
print("Slope ( $\beta_1$ ):", model.coef_)
print("Intercept ( $\beta_0$ ):", model.intercept_)
```

- **LinearRegression():** Creates a linear regression model.
- **model.fit:** Fits the model to the training data.
- **model.coef\_:** Extracts the slope ( $\beta_1$ ).
- **model.intercept\_:** Extracts the intercept ( $\beta_0$ ).

#### # Step 6: Make Predictions

```
y_pred = model.predict(X_test)
```

```
# Compare actual and predicted values
comparison = pd.DataFrame({'Actual': y_test.values, 'Predicted': y_pred})
print(comparison.head())
```

- **model.predict:** Generates predictions for the test set.
- **DataFrame:** Combines actual and predicted values for easy comparison.



#### Step 7: Evaluate the Model

```
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("\nModel Evaluation Metrics:")
print("Mean Squared Error (MSE):", mse)
print("R-squared (R2):", r2)
```

#### Step 8: Visualize the Results

```
# Scatter plot with regression line
plt.scatter(X_test, y_test, color='blue', label='Actual')
plt.plot(X_test, y_pred, color='red', label='Regression Line')
plt.xlabel("Avg. Area House Age")
plt.ylabel("Price")
plt.title("Linear Regression: Avg. Area House Age vs Price")
plt.legend() plt.show()
```

- **plt.scatter:** Creates a scatter plot of the actual data points.
- **plt.plot:** Adds the regression line to the plot.
- **xlabel, ylabel, title:** Adds labels and a title for better understanding.
- **legend:** Displays the legend.

#### 5. Summary and Insights

- **Slope ( $\beta_1$ ):** Indicates how much the dependent variable changes for a unit change in the independent variable.
- **Intercept ( $\beta_0$ ):** The predicted value of yyy when xxx is zero.
- **R-squared ( $R^2$ ):** Evaluates how well the model fits the data.
- **Visualization:** Confirms if the regression line captures the trend in the data.
- 

#### 6. Deliverables

1. A Jupyter Notebook with implementation.
2. Visualizations of the regression line and residuals.
3. Answers to lab questions.



UNIVERSITY OF ENGINEERING AND TECHNOLOGY, TAXILA

**FACULTY OF TELECOMMUNICATION AND INFORMATION  
ENGINEERING**

**COMPUTER ENGINEERING DEPARTMENT**

**Lab Questions**

1. What does the slope ( $\beta_1$ ) indicate for this dataset?
2. How well does the model predict  $y$  based on the  $R^2$  value?
3. Are there any patterns in the residuals that violate linear regression assumptions?